

# CSC413/2516 Lecture 10: Generative Models & Reinforcement Learning

Jimmy Ba and Bo Wang

# Logistics

Some administrative stuff:

- PA 4 (**most interesting**) is out! (Due April 1st, not a joke!)
- HW 4 (**most 'mathy'**) will be out in March 25th, and due April 08.
- Final Project is due April 12th! (likely to be extended, stay tuned!)

# Overview

Quiz: Which face image is fake?



A



B



C

# Overview

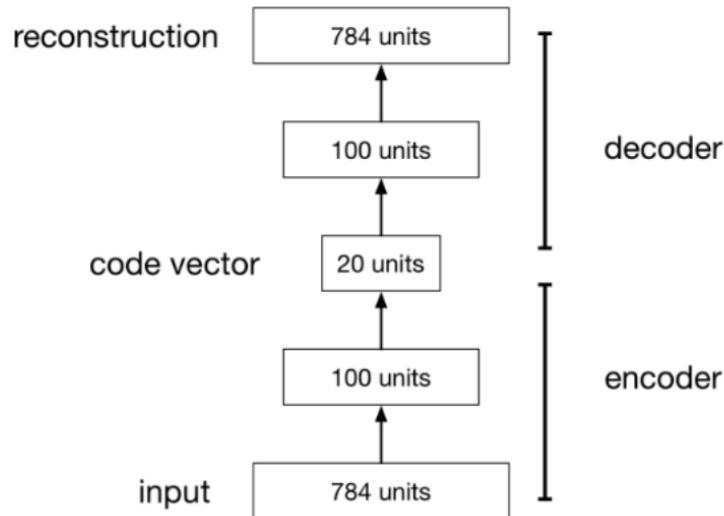
Four modern approaches to generative modeling:

- Autoregressive models (Lectures 3, 7, and 8)
- Generative adversarial networks (last lecture)
- Reversible architectures (last lecture)
- **Variational autoencoders (this lecture)**

All four approaches have different pros and cons.

# Autoencoders

- An **autoencoder** is a feed-forward neural net whose job it is to take an input  $x$  and predict  $x$ .
- To make this non-trivial, we need to add a **bottleneck layer** whose dimension is much smaller than the input.



# Autoencoders

## Why autoencoders?


- Map high-dimensional data to two dimensions for visualization
- Compression (i.e. reducing the file size)
  - Note: this requires a VAE, not just an ordinary autoencoder.
- Learn abstract features in an unsupervised way so you can apply them to a supervised task
  - Unlabeled data can be much more plentiful than labeled data
- Learn a semantically meaningful representation where you can, e.g., interpolate between different images.

## Principal Component Analysis (optional)

- The simplest kind of autoencoder has one hidden layer, linear activations, and squared error loss.

$$\mathcal{L}(\mathbf{x}, \tilde{\mathbf{x}}) = \|\mathbf{x} - \tilde{\mathbf{x}}\|^2$$

- This network computes  $\tilde{\mathbf{x}} = \mathbf{U}\mathbf{V}\mathbf{x}$ , which is a linear function.
- If  $K \geq D$ , we can choose  $\mathbf{U}$  and  $\mathbf{V}$  such that  $\mathbf{UV}$  is the identity. This isn't very interesting.
- But suppose  $K < D$ :
  - $\mathbf{V}$  maps  $\mathbf{x}$  to a  $K$ -dimensional space, so it's doing dimensionality reduction.
  - The output must lie in a  $K$ -dimensional subspace, namely the column space of  $\mathbf{U}$ .



# Principal Component Analysis (optional)



- Review from CSC311: linear autoencoders with squared error loss are equivalent to Principal Component Analysis (PCA).
- Two equivalent formulations:
  - Find the subspace that minimizes the reconstruction error.
  - Find the subspace that maximizes the projected variance.
- The optimal subspace is spanned by the dominant eigenvectors of the empirical covariance matrix.



“Eigenfaces”


# Deep Autoencoders

- Deep nonlinear autoencoders learn to project the data, not onto a subspace, but onto a **nonlinear manifold**
- This manifold is the image of the decoder.
- This is a kind of **nonlinear dimensionality reduction**.




# Deep Autoencoders

- Nonlinear autoencoders can learn more powerful codes for a given dimensionality, compared with linear autoencoders (PCA)




# Deep Autoencoders

- Some limitations of autoencoders
  - They're not generative models, so they don't define a distribution
  - How to choose the latent dimension?




# Variational Auto-encoder (VAE)




**Decoder** learns the **generative process** given the sampled latent vectors.

Sampling process in the middle.



**Encoder** learns the distribution of latent space given the observations.



# Variational Auto-encoder (VAE)



$$\text{loss} = \|x - \hat{x}\|^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)] = \|x - d(z)\|^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)]$$

Source: <https://towardsdatascience.com/understanding-variational-autoencoders-vae-f70510919f73>

# Observation Model



Source: <https://iagt.m.pressbooks.com/chapter/story-platos-allegory-of-the-cave/>

## Observation Model


- Consider training a generator network with maximum likelihood.

$$p(\mathbf{x}) = \int p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) d\mathbf{z}$$

- One problem: if  $\mathbf{z}$  is low-dimensional and the decoder is deterministic, then  $p(\mathbf{x}) = 0$  almost everywhere!
  - The model only generates samples over a low-dimensional sub-manifold of  $\mathcal{X}$ .
- Solution: define a noisy observation model, e.g.


$$p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{x}; G_\theta(\mathbf{z}), \eta \mathbf{I}),$$

where  $G_\theta$  is the function computed by the decoder with parameters  $\theta$ .



# Observation Model

- At least  $p(\mathbf{x}) = \int p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) d\mathbf{z}$  is well-defined, but how can we compute it?
- Integration, according to XKCD:



# Observation Model

- At least  $p(\mathbf{x}) = \int p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) d\mathbf{z}$  is well-defined, but how can we compute it?
  - The decoder function  $G_\theta(\mathbf{z})$  is very complicated, so there's no hope of finding a closed form.
- Instead, we will try to maximize a lower bound on  $\log p(\mathbf{x})$ .
  - The math is essentially the same as in the EM algorithm from CSC411.

# Variational Inference

- We obtain the lower bound using **Jensen's Inequality**: for a convex function  $h$  of a random variable  $X$ ,



$$\mathbb{E}[h(X)] \geq h(\mathbb{E}[X])$$

Therefore, if  $h$  is **concave** (i.e.  $-h$  is convex),

$$\mathbb{E}[h(X)] \leq h(\mathbb{E}[X])$$

- The function  $\log z$  is concave.  
Therefore,

$$\mathbb{E}[\log X] \leq \log \mathbb{E}[X]$$



# Variational Inference

- Suppose we have some distribution  $q(\mathbf{z})$ . (We'll see later where this comes from.)
- We use Jensen's Inequality to obtain the lower bound.

Heads-up: You will show the full calculation in HW4.

$$\begin{aligned}\log p(\mathbf{x}) &= \log \int p(\mathbf{z}) p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &= \log \int q(\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &\geq \int q(\mathbf{z}) \log \left[ \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) \right] d\mathbf{z} \quad (\text{Jensen's Inequality}) \\ &= \mathbb{E}_q \left[ \log \frac{p(\mathbf{z})}{q(\mathbf{z})} \right] + \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})]\end{aligned}$$

- We'll look at these two terms in turn.

# Variational Inference

- The first term we'll look at is  $\mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})]$
- Since we assumed a Gaussian observation model,

$$\begin{aligned}\log p(\mathbf{x}|\mathbf{z}) &= \log \mathcal{N}(\mathbf{x}; G_{\theta}(\mathbf{z}), \eta \mathbf{I}) \\ &= \log \left[ \frac{1}{(2\pi\eta)^{D/2}} \exp \left( -\frac{1}{2\eta} \|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2 \right) \right] \\ &= -\frac{1}{2\eta} \|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2 + \text{const}\end{aligned}$$

- So this term is the expected squared error in reconstructing  $\mathbf{x}$  from  $\mathbf{z}$ .  
We call it the **reconstruction term**.

# Variational Inference

- The second term is  $\mathbb{E}_q \left[ \log \frac{p(\mathbf{z})}{q(\mathbf{z})} \right]$ .
- This is just  $-D_{KL}(q(\mathbf{z}) \| p(\mathbf{z}))$ , where  $D_{KL}$  is the Kullback-Leibler (KL) divergence

$$D_{KL}(q(\mathbf{z}) \| p(\mathbf{z})) \triangleq \mathbb{E}_q \left[ \log \frac{q(\mathbf{z})}{p(\mathbf{z})} \right]$$

- KL divergence is a widely used measure of distance between probability distributions, though it doesn't satisfy the axioms to be a distance metric.
- More details in tutorial.
- Typically,  $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ . Hence, the KL term encourages  $q$  to be close to  $\mathcal{N}(\mathbf{0}, \mathbf{I})$ .

# Variational Inference

- Hence, we're trying to maximize the **variational lower bound**, or **variational free energy**:

$$\log p(\mathbf{x}) \geq \mathcal{F}(\boldsymbol{\theta}, q) = \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q\|p).$$

- The term “variational” is a historical accident: “variational inference” used to be done using variational calculus, but this isn’t how we train VAEs.
- We’d like to choose  $q$  to make the bound as tight as possible.
- It’s possible to show that the gap is given by:

$$\log p(\mathbf{x}) - \mathcal{F}(\boldsymbol{\theta}, q) = D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z}|\mathbf{x})).$$

Therefore, we’d like  $q$  to be as close as possible to the posterior distribution  $p(\mathbf{z}|\mathbf{x})$ .

- Let's think about the role of each of the two terms.
- The reconstruction term

$$\mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})] = -\frac{1}{2\sigma^2}\mathbb{E}_q[\|\mathbf{x} - G_\theta(\mathbf{z})\|^2] + \text{const}$$

is minimized when  $q$  is a **point mass** on

$$\mathbf{z}_* = \arg \min_{\mathbf{z}} \|\mathbf{x} - G_\theta(\mathbf{z})\|^2.$$

- But a point mass would have infinite KL divergence. (Exercise: check this.) So the KL term forces  $q$  to be more spread out.

# Reparameterization Trick

- To fit  $q$ , let's assign it a parametric form, in particular a Gaussian distribution:  $q(\mathbf{z}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$ , where  $\boldsymbol{\mu} = (\mu_1, \dots, \mu_K)$  and  $\boldsymbol{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_K^2)$ .
- In general, it's hard to differentiate through an expectation. But for Gaussian  $q$ , we can apply the **reparameterization trick**:

$$z_i = \mu_i + \sigma_i \epsilon_i,$$

where  $\epsilon_i \sim \mathcal{N}(0, 1)$ .


- Hence,

$$\overline{\mu_i} = \overline{z_i} \quad \overline{\sigma_i} = \overline{z_i} \epsilon_i.$$

- This is exactly analogous to how we derived the backprop rules for dropout

# Reparameterization Trick

Original form




: Deterministic node




: Random node

# Reparameterization Trick

Original form



Reparameterised form



: Deterministic node



: Random node

[Kingma, 2013]

[Bengio, 2013]

[Kingma and Welling 2014]


[Rezende et al 2014]

# Amortization

- This suggests one strategy for learning the decoder. For each training example,
  - ① Fit  $q$  to approximate the posterior for the current  $\mathbf{x}$  by doing many steps of gradient ascent on  $\mathcal{F}$ .
  - ② Update the decoder parameters  $\theta$  with gradient ascent on  $\mathcal{F}$ .
- **Problem:** this requires an expensive iterative procedure for every training example, so it will take a long time to process the whole training set.


# Amortization

- **Idea:** amortize the cost of inference by learning an inference network which predicts  $(\mu, \Sigma)$  as a function of  $\mathbf{x}$ .
- The outputs of the inference net are  $\mu$  and  $\log \sigma$ . (The log representation ensures  $\sigma > 0$ .)
- If  $\sigma \approx 0$ , then this network essentially computes  $\mathbf{z}$  deterministically, by way of  $\mu$ .
  - But the KL term encourages  $\sigma > 0$ , so in general  $\mathbf{z}$  will be noisy.
- The notation  $q(\mathbf{z}|\mathbf{x})$  emphasizes that  $q$  depends on  $\mathbf{x}$ , even though it's not actually a conditional distribution.




# Amortization

- Combining this with the decoder network, we see the structure closely resembles an ordinary autoencoder. The inference net is like an encoder.
- Hence, this architecture is known as a **variational autoencoder (VAE)**.
- The parameters of both the encoder and decoder networks are updated using a single pass of ordinary backprop.
  - The reconstruction term corresponds to squared error  $\|\mathbf{x} - \tilde{\mathbf{x}}\|^2$ , like in an ordinary VAE.
  - The KL term regularizes the representation by encouraging  $\mathbf{z}$  to be more stochastic.



# VAE - Summary



Push  $\mathbf{x}$  through encoder

$$\epsilon \sim \mathcal{N}(0, 1)$$

Sample noise

$$\mathbf{z} = \epsilon \sigma_x + \mu_x$$

Reparameterize

$$\mathbf{x}_r = p_\theta(\mathbf{x} \mid \mathbf{z})$$

Push  $\mathbf{z}$  through decoder

$$\text{recon. loss} = \text{MSE}(\mathbf{x}, \mathbf{x}_r)$$

Compute reconstruction loss

$$\text{var. loss} = -\text{KL}[\mathcal{N}(\mu_x, \sigma_x) \parallel \mathcal{N}(0, I)]$$

Compute variational loss

$$L = \text{recon. loss} + \text{var. loss}$$

Combine losses

## VAEs vs. Other Generative Models

- In short, a VAE is like an autoencoder, except that it's also a generative model (defines a distribution  $p(\mathbf{x})$ ).
- Unlike autoregressive models, generation only requires one forward pass.
- Unlike reversible models, we can fit a low-dimensional latent representation. We'll see we can do interesting things with this...



# Latent Space Interpolations


- You can often get interesting results by interpolating between two vectors in the latent space:



Ha and Eck, "A neural representation of sketch drawings"

# Latent Space Interpolations

- You can often get interesting results by interpolating between two vectors in the latent space:



<https://arxiv.org/pdf/1610.00291.pdf>

# Latent Space Interpolations

- Latent space interpolation of music:

<https://magenta.tensorflow.org/music-vae>

After the break

After the break: **Reinforcement Learning: Policy Gradient**


# Alpha-Go Trailer --- This is it, folks!



# Overview

- Most of this course was about supervised learning, plus a little unsupervised learning.
- Reinforcement learning:
  - Middle ground between supervised and unsupervised learning
  - An agent acts in an environment and receives a reward signal.
- Today: policy gradient (directly do SGD over a stochastic policy using trial-and-error)
- Next lecture: combine policies and Q-learning

# Overview



Source: <https://perfectial.com/blog/reinforcement-learning-applications/>

# Overview

How does AI take over the world? Three Steps!


**SUPERVISED  
LEARNING**

**UNSUPERVISED  
LEARNING**

**REINFORCEMENT  
LEARNING**



# Reinforcement learning



- An **agent** interacts with an **environment** (e.g. game of Breakout)
- In each time step  $t$ ,
  - the agent receives **observations** (e.g. pixels) which give it information about the **state**  $s_t$  (e.g. positions of the ball and paddle)
  - the agent picks an **action**  $a_t$  (e.g. keystrokes) which affects the state
- The agent periodically receives a **reward**  $r(s_t, a_t)$ , which depends on the state and action (e.g. points)
- The agent wants to learn a **policy**  $\pi_\theta(a_t | s_t)$ 
  - Distribution over actions depending on the current state and parameters  $\theta$

# Reinforcement learning

## Cart-Pole Problem



**Objective:** Balance a pole on top of a movable cart

**State:** angle, angular speed, position, horizontal velocity

**Action:** horizontal force applied on the cart

**Reward:** 1 at each time step if the pole is upright

Source: Fei-Fei Li, Justin Johnson, Serena Yeung, **cs231n Stanford**

# Reinforcement learning

## Robot Locomotion



**Objective:** Make the robot move forward

**State:** Angle and position of the joints


**Action:** Torques applied on joints

**Reward:** 1 at each time step upright + forward movement

Source: Fei-Fei Li, Justin Johnson, Serena Yeung, **cs231n Stanford**

# Reinforcement learning

## Go



**Objective:** Win the game!

**State:** Position of all pieces

**Action:** Where to put the next piece down

**Reward:** 1 if win at the end of the game, 0 otherwise

Source: Fei-Fei Li, Justin Johnson, Serena Yeung, [cs231n Stanford](#)


# Markov Decision Processes

- The environment is represented as a **Markov decision process**  $\mathcal{M}$ .
- Markov assumption: all relevant information is encapsulated in the current state; i.e. the policy, reward, and transitions are all independent of past states given the current state
- Components of an MDP:
  - initial state distribution  $p(\mathbf{s}_0)$
  - policy  $\pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$
  - transition distribution  $p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$
  - reward function  $r(\mathbf{s}_t, \mathbf{a}_t)$
- Assume a **fully observable** environment, i.e.  $\mathbf{s}_t$  can be observed directly
- **Rollout**, or **trajectory**  $\tau = (\mathbf{s}_0, \mathbf{a}_0, \mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_T, \mathbf{a}_T)$
- Probability of a rollout

$$p(\tau) = p(\mathbf{s}_0) \pi_\theta(\mathbf{a}_0 | \mathbf{s}_0) p(\mathbf{s}_1 | \mathbf{s}_0, \mathbf{a}_0) \cdots p(\mathbf{s}_T | \mathbf{s}_{T-1}, \mathbf{a}_{T-1}) \pi_\theta(\mathbf{a}_T | \mathbf{s}_T)$$

# Markov Decision Processes



Continuous control in simulation, e.g. teaching an ant to walk



- State: positions, angles, and velocities of the joints
- Actions: apply forces to the joints
- Reward: distance from starting point
- Policy: output of an ordinary MLP, using the state as input
- More environments: <https://gym.openai.com/envs/#mujoco>

# Markov Decision Processes

- **Return** for a rollout:  $r(\tau) = \sum_{t=0}^T r(s_t, a_t)$ 
  - Note: we're considering a finite **horizon**  $T$ , or number of time steps; we'll consider the infinite horizon case later.
- Goal: maximize the expected return,  $R = \mathbb{E}_{p(\tau)}[r(\tau)]$
- The expectation is over both the environment's dynamics and the policy, but we only have control over the policy.
- The stochastic policy is important, since it makes  $R$  a continuous function of the policy parameters.
  - Reward functions are often discontinuous, as are the dynamics (e.g. collisions)



## REINFORCE (Policy Gradient)

- REINFORCE is an elegant algorithm for maximizing the expected return  $R = \mathbb{E}_{p(\tau)} [r(\tau)]$ .
- Intuition: trial and error
  - Sample a rollout  $\tau$ . If you get a high reward, try to make it more likely.  
If you get a low reward, try to make it less likely.
- Interestingly, this can be seen as stochastic gradient ascent on  $R$ .

# REINFORCE

- Recall the derivative formula for log:

Important Trick!

$$\frac{\partial}{\partial \theta} \log p(\tau) = \frac{\frac{\partial}{\partial \theta} p(\tau)}{p(\tau)} \implies \frac{\partial}{\partial \theta} p(\tau) = p(\tau) \frac{\partial}{\partial \theta} \log p(\tau)$$

- Gradient of the expected return:

$$\begin{aligned}\frac{\partial}{\partial \theta} \mathbb{E}_{p(\tau)} [r(\tau)] &= \frac{\partial}{\partial \theta} \sum_{\tau} r(\tau) p(\tau) \\&= \sum_{\tau} r(\tau) \frac{\partial}{\partial \theta} p(\tau) \\&= \sum_{\tau} r(\tau) p(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \\&= \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right]\end{aligned}$$

- Compute stochastic estimates of this expectation by sampling rollouts.

# REINFORCE

- For reference:

$$\frac{\partial}{\partial \theta} \mathbb{E}_{p(\tau)} [r(\tau)] = \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right]$$

- If you get a large reward, make the rollout more likely. If you get a small reward, make it less likely.
- Unpacking the REINFORCE gradient:

$$\begin{aligned}\frac{\partial}{\partial \theta} \log p(\tau) &= \frac{\partial}{\partial \theta} \log \left[ p(\mathbf{s}_0) \prod_{t=0}^T \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) \prod_{t=1}^T p(\mathbf{s}_t | \mathbf{s}_{t-1}, \mathbf{a}_{t-1}) \right] \\ &= \frac{\partial}{\partial \theta} \log \prod_{t=0}^T \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) \\ &= \sum_{t=0}^T \frac{\partial}{\partial \theta} \log \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)\end{aligned}$$

- Hence, it tries to make *all* the actions more likely or less likely, depending on the reward. I.e., it doesn't do **credit assignment**.
  - This is a topic for next lecture.

# REINFORCE

Repeat forever:

Sample a rollout  $\tau = (\mathbf{s}_0, \mathbf{a}_0, \mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_T, \mathbf{a}_T)$

$$r(\tau) \leftarrow \sum_{k=0}^T r(\mathbf{s}_k, \mathbf{a}_k)$$

For  $t = 0, \dots, T$ :

$$\theta \leftarrow \theta + \alpha r(\tau) \frac{\partial}{\partial \theta} \log \pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$$

- Observation: actions should only be reinforced based on future rewards, since they can't possibly influence past rewards.
- You can show that this still gives unbiased gradient estimates.

Repeat forever:

Sample a rollout  $\tau = (\mathbf{s}_0, \mathbf{a}_0, \mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_T, \mathbf{a}_T)$


For  $t = 0, \dots, T$ :

$$r_t(\tau) \leftarrow \sum_{k=t}^T r(\mathbf{s}_k, \mathbf{a}_k)$$


$$\theta \leftarrow \theta + \alpha r_t(\tau) \frac{\partial}{\partial \theta} \log \pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$$

- A classification task under RL formulation
  - one time step
  - state  $x$ : an image
  - action  $a$ : a digit class
  - reward  $r(x, a)$ : 1 if correct, 0 if wrong
  - policy  $\pi(a|x)$ : a distribution over categories
    - Compute using an MLP with softmax outputs – this is a **policy network**

# RL for Classification



Non-differentiable: OK



Discontinuous: not OK

- Original solution: use a surrogate loss function, e.g. logistic-cross-entropy
- RL formulation: in each episode, the agent is shown an image, guesses a digit class, and receives a reward of 1 if it's right or 0 if it's wrong
- We'd never actually do it this way, but it will give us an interesting comparison with backprop.

## RL for Classification

- Let  $z_k$  denote the logits,  $y_k$  denote the softmax output,  $t$  the integer target, and  $t_k$  the target one-hot representation.
- To apply REINFORCE, we sample  $\mathbf{a} \sim \pi_{\theta}(\cdot | \mathbf{x})$  and apply:

$$\begin{aligned}\theta &\leftarrow \theta + \alpha r(\mathbf{a}, t) \frac{\partial}{\partial \theta} \log \pi_{\theta}(\mathbf{a} | \mathbf{x}) \\&= \theta + \alpha r(\mathbf{a}, t) \frac{\partial}{\partial \theta} \log y_a \\&= \theta + \alpha r(\mathbf{a}, t) \sum_k (a_k - y_k) \frac{\partial}{\partial \theta} z_k\end{aligned}$$

- Compare with the logistic regression SGD update:

$$\begin{aligned}\theta &\leftarrow \theta + \alpha \frac{\partial}{\partial \theta} \log y_t \\&\leftarrow \theta + \alpha \sum_k (t_k - y_k) \frac{\partial}{\partial \theta} z_k\end{aligned}$$

## Reward Baselines

- For reference:

$$\theta \leftarrow \theta + \alpha r(\mathbf{a}, \mathbf{t}) \frac{\partial}{\partial \theta} \log \pi_{\theta}(\mathbf{a} | \mathbf{x})$$

- Clearly, we can add a constant offset to the reward, and we get an equivalent optimization problem.
- Behavior if  $r = 0$  for wrong answers and  $r = 1$  for correct answers
  - wrong: do nothing
  - correct: make the action more likely
- If  $r = 10$  for wrong answers and  $r = 11$  for correct answers
  - wrong: make the action more likely
  - correct: make the action more likely (slightly stronger)
- If  $r = -10$  for wrong answers and  $r = -9$  for correct answers
  - wrong: make the action less likely
  - correct: make the action less likely (slightly weaker)

# Reward Baselines

- Problem: the REINFORCE update depends on arbitrary constant factors added to the reward.
- Observation: we can subtract a **baseline**  $b$  from the reward without biasing the gradient.

$$\begin{aligned}\mathbb{E}_{p(\tau)} \left[ (r(\tau) - b) \frac{\partial}{\partial \theta} \log p(\tau) \right] &= \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right] - b \mathbb{E}_{p(\tau)} \left[ \frac{\partial}{\partial \theta} \log p(\tau) \right] \\ &= \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right] - b \sum_{\tau} p(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \\ &= \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right] - b \sum_{\tau} \frac{\partial}{\partial \theta} p(\tau) \\ &= \mathbb{E}_{p(\tau)} \left[ r(\tau) \frac{\partial}{\partial \theta} \log p(\tau) \right] - 0 \quad \text{Hint: } \sum_{\tau} p_{\theta}(\tau) = 1\end{aligned}$$

Heads-up: You will show how  $b$  can be selected to reduce variance in HW4.

- We'd like to pick a baseline such that good rewards are positive and bad ones are negative.
- $\mathbb{E}[r(\tau)]$  is a good choice of baseline, but we can't always compute it easily. There's lots of research on trying to approximate it.

## More Tricks


- We left out some more tricks that can make policy gradients work a lot better.
  - Natural policy gradient corrects for the geometry of the space of policies, preventing the policy from changing too quickly.
  - Rather than use the actual return, evaluate actions based on estimates of future returns. This is a class of methods known as actor-critic, which we'll touch upon next lecture.
- Trust region policy optimization (TRPO) and proximal policy optimization (PPO) are modern policy gradient algorithms which are very effective for continuous control problems.

# Discussion

- What's so great about backprop and gradient descent?
  - Backprop does credit assignment – it tells you exactly which activations and parameters should be adjusted upwards or downwards to decrease the loss on some training example.
  - REINFORCE doesn't do credit assignment. If a rollout happens to be good, all the actions get reinforced, even if some of them were bad.
  - Reinforcing all the actions as a group leads to random walk behavior.

# Discussion

- Why policy gradient?
  - Can handle discontinuous cost functions
  - Don't need an explicit model of the environment, i.e. rewards and dynamics are treated as black boxes
    - Policy gradient is an example of **model-free reinforcement learning**, since the agent doesn't try to fit a model of the environment
    - Almost everyone thinks model-based approaches are needed for AI, but nobody has a clue how to get it to work

A photograph of a dirt road winding through a forest. The road is light brown and appears slightly damp. On either side are tall, dark evergreen trees, with some bare deciduous branches visible. The sky is overcast and grey.

**What I'm trying to do is to  
maximise the probability  
of the future being better.**

Elon Musk